

# **EE3801 Data Engineering Principles**

**Parallel Data Processing**

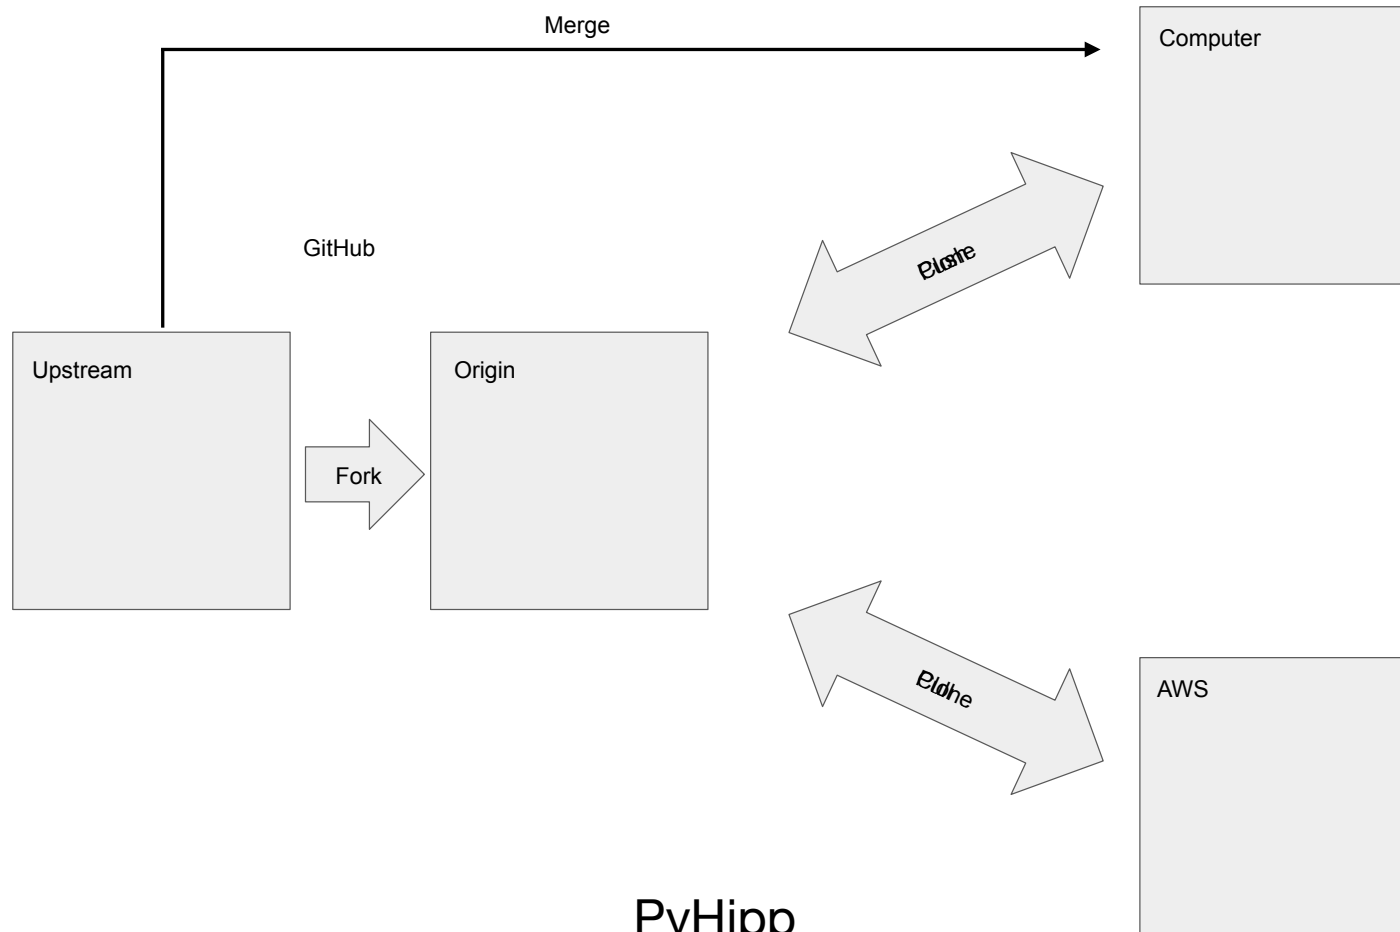
# Parallel Data Processing

## Parallel Processing

- Lab 5 techniques
- Parallelizing jobs
- Increasing memory per job
- Terminating cluster head nodes automatically
- Slurm job dependencies

# GitHub Repositories

## Forked Repositories



# GitHub Repositories

## Forked Repositories

On AWS

Check that the upstream repository is configured correctly:

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git remote -v
```

which should return:

```
origin      https://github.com/yourusername/PyHipp.git (fetch)
origin      https://github.com/yourusername/PyHipp.git (push)
upstream    https://github.com/shihchengyen/PyHipp.git (fetch)
upstream    https://github.com/shihchengyen/PyHipp.git (push)
```

If the upstream repository is not set, you can add it by doing:

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git remote add upstream https://github.com/shihc
```

On AWS

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git fetch upstream
```

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git checkout main
```

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git merge upstream/main
```

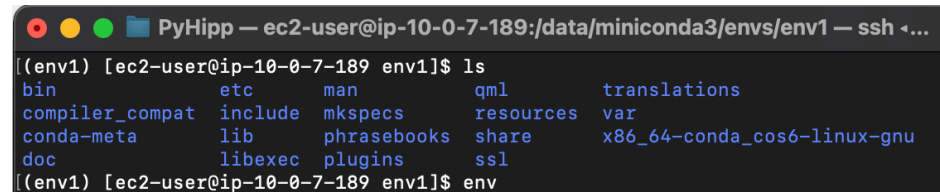
This should add a file named "ec2snapshot.sh". You can push the merged changes to your forked repository by doing:

```
(env1) [ec2-user@ip-10-0-5-43 PyHipp] $ git push
```

# Optimizing Parallel Processing

## Lab 5 Techniques (Git and Conda)

- GitHub command line (git clone)
- Conda environments (e.g. /data/miniconda3/envs/env1)



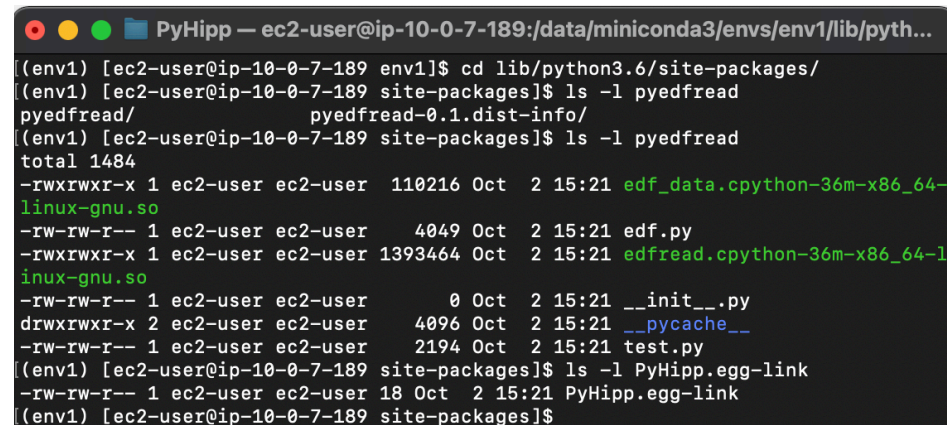
```
PyHipp — ec2-user@ip-10-0-7-189:/data/miniconda3/envs/env1 — ssh ◀...  
[(env1) [ec2-user@ip-10-0-7-189 env1]$ ls  
bin          etc          man          qml          translations  
compiler_compat  include    mkspecs     resources    var  
conda-meta    lib         phrasebooks share         x86_64-conda_cos6-linux-gnu  
doc           libexec     plugins      ssl  
[(env1) [ec2-user@ip-10-0-7-189 env1]$ env
```

- Software is installed in /data/miniconda3/envs/env1/lib/python3.8/site-packages
- Python checks this directory when functions are called

# Optimizing Parallel Processing

## Lab 5 Techniques (Conda)

- `pip install .`
- copies code to `/data/miniconda3/envs/env1/lib/python3.8/site-packages`, e.g. `pyedfread`



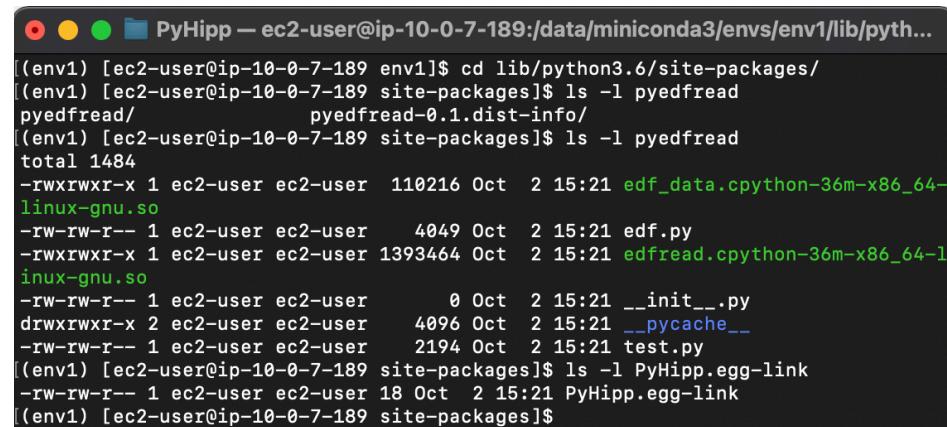
```
PyHipp — ec2-user@ip-10-0-7-189:/data/miniconda3/envs/env1/lib/pyth...
(env1) [ec2-user@ip-10-0-7-189 env1]$ cd lib/python3.6/site-packages/
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l pyedfread
pyedfread/
pyedfread-0.1.dist-info/
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l pyedfread
total 1484
-rwxrwxr-x 1 ec2-user ec2-user 110216 Oct  2 15:21 edf_data.cpython-36m-x86_64-
linux-gnu.so
-rw-rw-r-- 1 ec2-user ec2-user  4049 Oct  2 15:21 edf.py
-rwxrwxr-x 1 ec2-user ec2-user 1393464 Oct  2 15:21 edfread.cpython-36m-x86_64-1
inux-gnu.so
-rw-rw-r-- 1 ec2-user ec2-user      0 Oct  2 15:21 __init__.py
drwxrwxr-x 2 ec2-user ec2-user  4096 Oct  2 15:21 __pycache__
-rw-rw-r-- 1 ec2-user ec2-user  2194 Oct  2 15:21 test.py
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l PyHipp.egg-link
-rw-rw-r-- 1 ec2-user ec2-user 18 Oct  2 15:21 PyHipp.egg-link
(env1) [ec2-user@ip-10-0-7-189 site-packages]$
```

- Each time the code changes, you have to do:
  - `git pull` ← (updates `/data/src/pyedfread`)
  - `pip install .` ← (copies to `/data/miniconda3/envs/env1/lib/python3.8/site-packages`)
- You can then start using the new function(s)

# Optimizing Parallel Processing

## Lab 5 Techniques (Bash)

- `pip install -e .`
- Links to `/data/miniconda3/envs/env1/lib/python3.8/site-packages`, e.g. PyHipp



```
PyHipp — ec2-user@ip-10-0-7-189:/data/miniconda3/envs/env1/lib/pyth...
(env1) [ec2-user@ip-10-0-7-189 env1]$ cd lib/python3.6/site-packages/
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l pyedfread
pyedfread/
pyedfread-0.1.dist-info/
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l pyedfread
total 1484
-rwxrwxr-x 1 ec2-user ec2-user 110216 Oct  2 15:21 edf_data.cpython-36m-x86_64-
linux-gnu.so
-rw-rw-r-- 1 ec2-user ec2-user  4049 Oct  2 15:21 edf.py
-rwxrwxr-x 1 ec2-user ec2-user 1393464 Oct  2 15:21 edfread.cpython-36m-x86_64-l
inux-gnu.so
-rw-rw-r-- 1 ec2-user ec2-user      0 Oct  2 15:21 __init__.py
drwxrwxr-x 2 ec2-user ec2-user  4096 Oct  2 15:21 __pycache__
-rw-rw-r-- 1 ec2-user ec2-user  2194 Oct  2 15:21 test.py
(env1) [ec2-user@ip-10-0-7-189 site-packages]$ ls -l PyHipp.egg-link
-rw-rw-r-- 1 ec2-user ec2-user 18 Oct  2 15:21 PyHipp.egg-link
(env1) [ec2-user@ip-10-0-7-189 site-packages]$
```

- Each time the code changes, all you have to do is:
  - `git pull` ← (updates `/data/src/pyedfread`)
- You can then start using the new function(s)

# Optimizing Parallel Processing

## Lab 5 Techniques (Compute Nodes)

**Compute Node**

Last Digit	Type	Last Digit	Type	Last Digit	Type
0	m5.4xlarge	4	r5n.2xlarge	8	r5.2xlarge
1	z1d.2xlarge	5	r5b.2xlarge	9	r5a.2xlarge
2	m5a.4xlarge	6	r5d.2xlarge		
3	r5dn.2xlarge	7	r5ad.2xlarge		

Instance Size	vCPU	Memory (GiB)	Instance Storage (GB)	Network Bandwidth (Gbps)***	EBS Bandwidth (Mbps)
m5.large	2	8	EBS-Only	Up to 10	Up to 4,750
m5.xlarge	4	16	EBS-Only	Up to 10	Up to 4,750
m5.2xlarge	8	32	EBS-Only	Up to 10	Up to 4,750
m5.4xlarge	16	64	EBS-Only	Up to 10	4,750

Instance	vCPU	Mem (GiB)	Networking Performance (Gbps)***	SSD Storage (GB)
z1d.large	2	16	Up to 10	1 x 75 NVMe SSD
z1d.xlarge	4	32	Up to 10	1 x 150 NVMe SSD
z1d.2xlarge	8	64	Up to 10	1 x 300 NVMe SSD

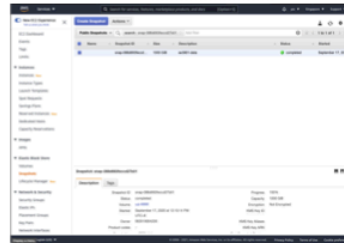


# Lab 3

## Snapshot

- Copy data snapshot

Click on the "Owned By Me" button to reveal the drop-down menu and select "Public Snapshots". Copy and paste the following snapshot id: snap-05c0a250a5fa4d56d, and then hit "Return". You should see a snapshot with the description "ee3801-2021-data" listed and selected.



Click on the "Actions" button, and select the Copy command. In the window that appears, replace the description with "data", make sure the "Encrypt this snapshot" option is NOT selected, and then click the "Copy" button. You should see a message that the snapshot is being copied.

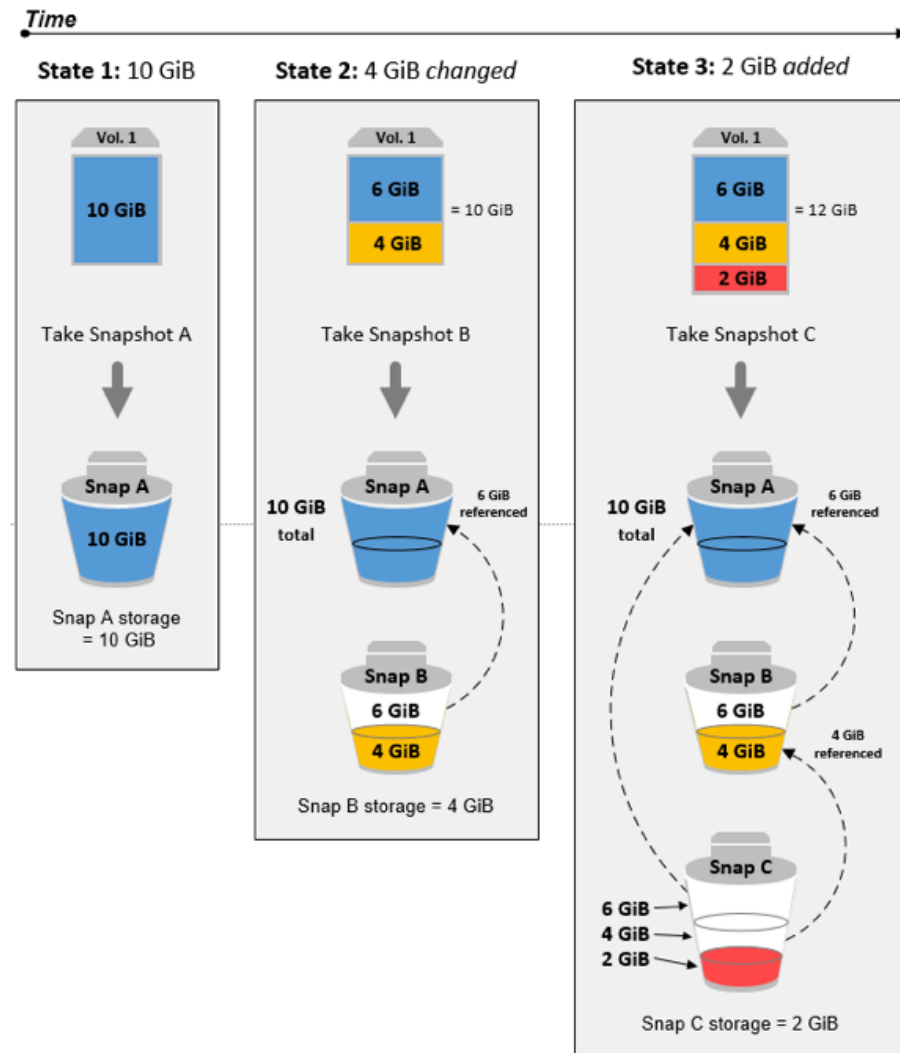
Click on the "Public Snapshots" button to select the "Owned By Me" option, and then click on the "x" icon in the search field to remove the snapshot id. You should now see a snapshot with a size of 1000 GiB with the description "data". If the snapshot is not selected, select it, and you should see more information for the snapshot shown in the panel at the bottom of the window. Move your cursor over the text "Snapshot ID", which will cause a "Copy to clipboard" icon to appear. Click on it to copy the Snapshot ID.

- Subsequent snapshots will be faster

# AWS ParallelCluster

## EBS Snapshots

- Incremental backups
- Only blocks that have changed after your most recent snapshot are saved
- Unchanged blocks are referenced to previous snapshot
- New data will add to the storage required
- Total storage for 3 snapshots is  $10 + 4 + 2 = 16$  GiB



# Optimizing Parallel Processing

## Lab 5 Techniques (Bash)

- ipython (easier to use than regular Python, e.g. ls, cd)
- tail -f file (to display text as they are added to the file)
- find (to find files and directories)
  - find . -name "\*.hkl"

```
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl"]  
./20181105/session01/data_raw6.hkl  
./20181105/session01/rplparallel_d41d.hkl  
./20181105/session01/unity_71bf.hkl  
(aws) Work27:dataeng shihcheng$
```

# Optimizing Parallel Processing

## Lab 5 Techniques (Bash)

- | (pipe output from one program to input of another program)

```
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl" ]
./20181105/session01/data_raw6.hkl
./20181105/session01/rplparallel_d41d.hkl
./20181105/session01/unity_71bf.hkl
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl" | grep unity ]
./20181105/session01/unity_71bf.hkl
(aws) Work27:dataeng shihcheng$ █
```

# Optimizing Parallel Processing

## Lab 5 Techniques (Bash)

- xargs (converts output from one program to arguments for another program)

```
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl"
./20181105/session01/data_raw6.hkl
./20181105/session01/rplparallel_d41d.hkl
./20181105/session01/unity_71bf.hkl
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl" | grep unity
./20181105/session01/unity_71bf.hkl
[(aws) Work27:dataeng shihcheng$ find . -name "*.hkl" | xargs ls -hl
-rwxrwxrwx 1 shihcheng staff 629M Sep 19 2021 ./20181105/session01/data_raw
6.hkl
-rwxrwxrwx 1 shihcheng staff 96K Sep 20 2021 ./20181105/session01/rplparal
lel_d41d.hkl
-rwxrwxrwx 1 shihcheng staff 11M Sep 20 2021 ./20181105/session01/unity_71
bf.hkl
[(aws) Work27:dataeng shihcheng$ ls -hl ./20181105/session01/data_raw6.hkl ./2018
1105/session01/rplparallel_d41d.hkl ./20181105/session01/unity_71bf.hkl
-rwxrwxrwx 1 shihcheng staff 629M Sep 19 2021 ./20181105/session01/data_raw
6.hkl
-rwxrwxrwx 1 shihcheng staff 96K Sep 20 2021 ./20181105/session01/rplparal
lel_d41d.hkl
-rwxrwxrwx 1 shihcheng staff 11M Sep 20 2021 ./20181105/session01/unity_71
bf.hkl
(aws) Work27:dataeng shihcheng$
```

# Optimizing Parallel Processing

## Lab 5 Techniques (Slurm)

- `srun --pty /bin/bash` (get shell running on compute node)
- `#SBATCH --time=24:00:00` # walltime (time limit for job)

# Parallel Data Processing

## Parallel Processing

- Lab 4 techniques
- Parallelizing jobs
- Increasing memory per job
- Terminating cluster head nodes automatically
- Slurm job dependencies

# Data Processing on AWS

## Slurm Script (Called from /data/picasso/20181105)

```
#!/bin/bash

# Submit this script with: sbatch <this-filename>

#SBATCH --time=24:00:00 # walltime
#SBATCH --ntasks=1 # number of processor cores (i.e. tasks)
#SBATCH --nodes=1 # number of nodes
#SBATCH -J "pipe" # job name

## /SBATCH -p general # partition (queue)
#SBATCH -o pipe-slurm.%N.%j.out # STDOUT
#SBATCH -e pipe-slurm.%N.%j.err # STDERR

# LOAD MODULES, INSERT CODE, AND RUN YOUR PROGRAMS HERE

python -u -c "import PyHipp as pyh; \
import DataProcessingTools as DPT; \
import os; \
import time; \
t0 = time.time(); \
print(time.localtime()); \
DPT.objects.processDirs(dirs=None, objtype=pyh.RPLParallel, saveLevel=1); \
DPT.objects.processDirs(dirs=None, objtype=pyh.RPLSplit, channel=[9, 31, 34, 56, 72, 93, 119, 120]); \
DPT.objects.processDirs(dirs=None, objtype=pyh.RPLLFP, saveLevel=1); \
DPT.objects.processDirs(dirs=None, objtype=pyh.RPLHighPass, saveLevel=1); \
DPT.objects.processDirs(dirs=None, objtype=pyh.Unity, saveLevel=1); \
pyh.EDFSplit(); \
os.chdir('session01'); \
pyh.aligning_objects(); \
pyh.raycast(1); \
DPT.objects.processDirs(level='channel', cmd='import PyHipp as pyh; from PyHipp import mountain_batch; mountain_batch.mountain_batch(); from PyHipp \
import export_mountain_cells; export_mountain_cells.export_mountain_cells();'); \
print(time.localtime()); \
print(time.time()-t0);"

aws sns publish --topic-arn arn:aws:sns:ap-southeast-1:123456789012:awsnotify --message "JobDone"
```



# Parallel Data Processing

## Parallel Processing

Serial Pipeline

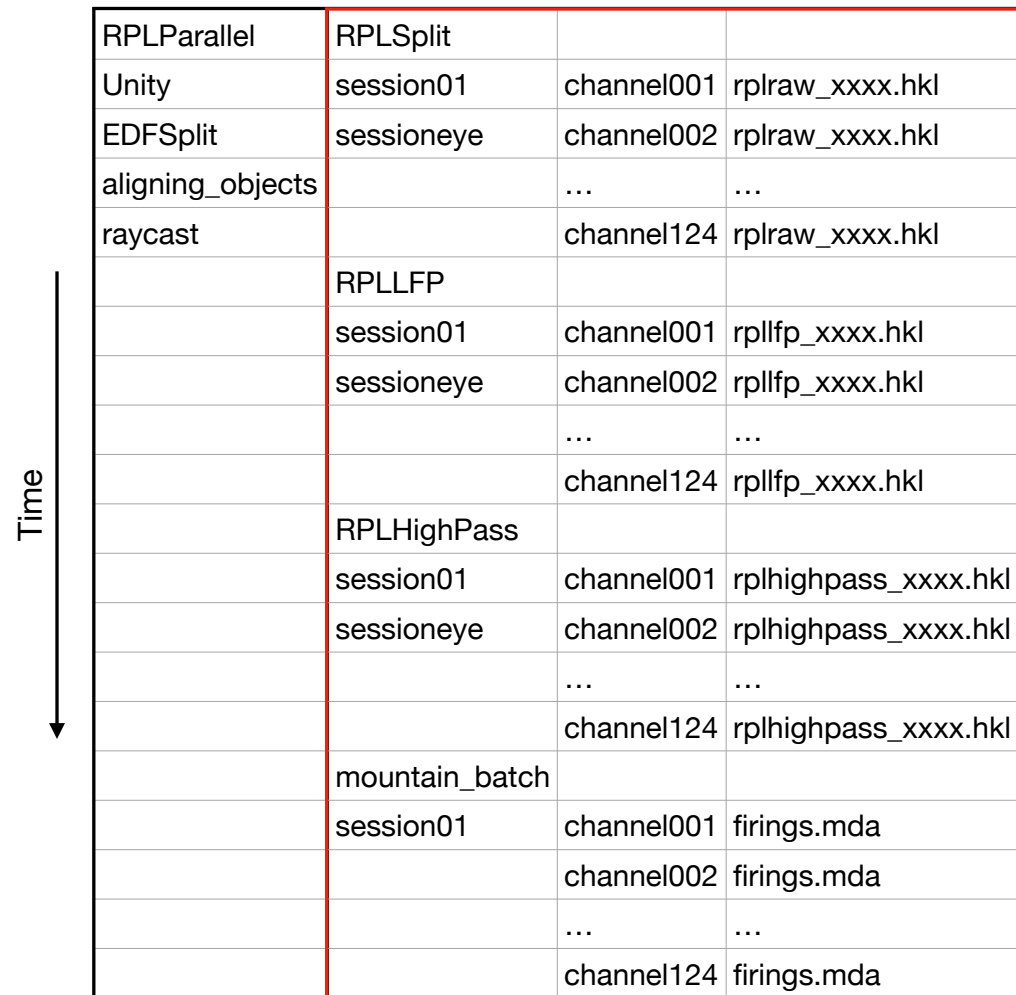
Dependencies	RPLParallel	RPLSplit	RPLLFP	RPLHighPass	Unity	EDFSplit
RPLParallel						
RPLSplit						
RPLLFP		√				
RPLHighPass		√				
Unity	√					
EDFSplit	√					
aligning_objects	√				√	√
raycast					√	√
mountain_batch				√		

Parallel Pipeline

RPLParallel	RPLSplit
Unity	RPLLFP
EDFSplit	RPLHighPass
aligning_objects	mountain_batch
raycast	

# Optimizing Parallel Processing

## Parallel Processing (2 jobs)



RPLParallel	RPLSplit		
Unity	session01	channel001	rplraw_xxxx.hkl
EDFSplit	sessioneye	channel002	rplraw_xxxx.hkl
aligning_objects		...	...
raycast		channel124	rplraw_xxxx.hkl
	RPLFP		
	session01	channel001	rplfp_xxxx.hkl
	sessioneye	channel002	rplfp_xxxx.hkl
		...	...
		channel124	rplfp_xxxx.hkl
	RPLHighPass		
	session01	channel001	rplhighpass_xxxx.hkl
	sessioneye	channel002	rplhighpass_xxxx.hkl
		...	...
		channel124	rplhighpass_xxxx.hkl
	mountain_batch		
	session01	channel001	firings.mda
		channel002	firings.mda
		...	...
		channel124	firings.mda

# Optimizing Parallel Processing

## Parallel Processing (5 jobs)

Time ↓

RPLParallel	RPLSplit		RPLSplit		RPLSplit		RPLSplit	
Unity	session01	channel001	session01	channel033	session01	channel065	session01	channel097
EDFSplit	sessioneye	channel002	sessioneye	channel034	sessioneye	channel066	sessioneye	channel098
aligning_objects		...		...		...		...
raycast		channel032		channel064		channel096		channel124
	RPLLFP		RPLLFP		RPLLFP		RPLLFP	
	session01	channel001	session01	channel033	session01	channel065	session01	channel097
	sessioneye	channel002	sessioneye	channel034	sessioneye	channel066	sessioneye	channel098
		...		...		...		...
		channel032		channel064		channel096		channel124
	RPLHighPass		RPLHighPass		RPLHighPass		RPLHighPass	
	session01	channel001	session01	channel033	session01	channel065	session01	channel097
	sessioneye	channel002	sessioneye	channel034	sessioneye	channel066	sessioneye	channel098
		...		...		...		...
		channel032		channel064		channel096		channel124
	mountain_batch		mountain_batch		mountain_batch		mountain_batch	
	session01	channel001	session01	channel033	session01	channel065	session01	channel097
		channel002		channel034		channel066		channel098
		...		...		...		...
		channel032		channel064		channel096		channel124

# Parallel Data Processing

## Parallel Processing

- Lab 4 techniques
- Parallelizing jobs
- Increasing memory per job
- Terminating cluster head nodes automatically
- Slurm job dependencies

# Optimizing Parallel Processing

## Parallel Processing (5 jobs)

Increasing memory per job

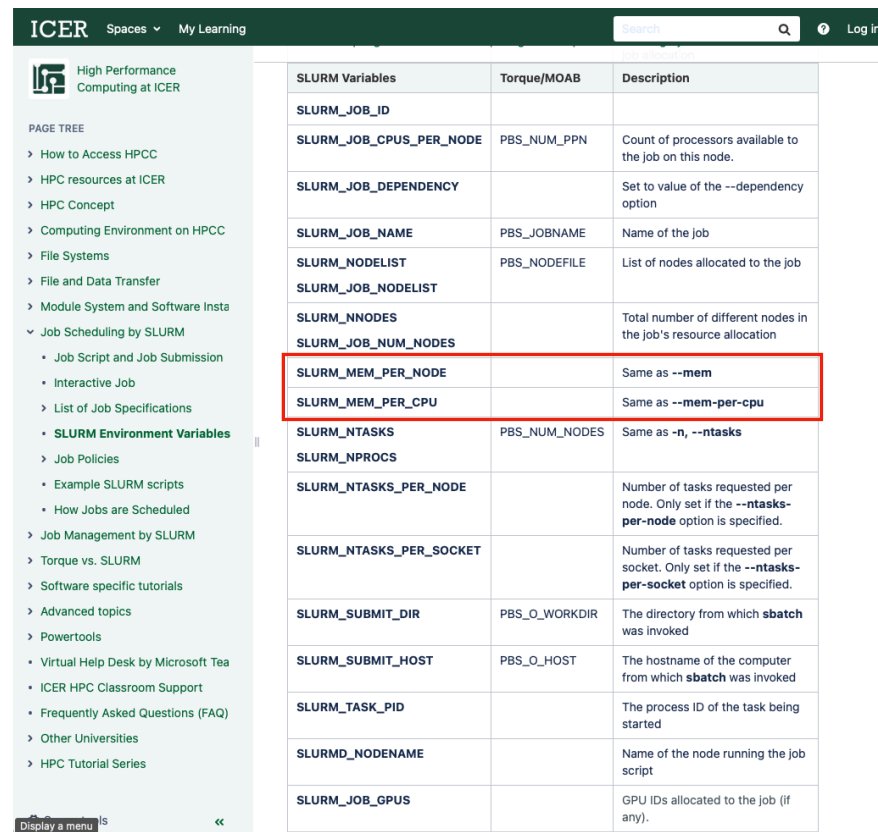
JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)	
2	queue1	rplp1	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	← 64 GB RAM
3	queue1	rs1	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
4	queue1	rs2	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
5	queue1	rs3	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
6	queue1	rs4	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	

↑  
40 GB RAM

# Optimizing Parallel Processing

## Parallel Processing (5 jobs)

### Increasing memory per job



SLURM Variables	Torque/MOAB	Description
SLURM_JOB_ID		
SLURM_JOB_CPUS_PER_NODE	PBS_NUM_PPN	Count of processors available to the job on this node.
SLURM_JOB_DEPENDENCY		Set to value of the --dependency option
SLURM_JOB_NAME	PBS_JOBNAME	Name of the job
SLURM_NODELIST	PBS_NODEFILE	List of nodes allocated to the job
SLURM_JOB_NODELIST		
SLURM_NNODES		Total number of different nodes in the job's resource allocation
SLURM_JOB_NUM_NODES		
SLURM_MEM_PER_NODE		Same as --mem
SLURM_MEM_PER_CPU		Same as --mem-per-cpu
SLURM_NTASKS	PBS_NUM_NODES	Same as -n, --ntasks
SLURM_NPROCS		
SLURM_NTASKS_PER_NODE		Number of tasks requested per node. Only set if the --ntasks-per-node option is specified.
SLURM_NTASKS_PER_SOCKET		Number of tasks requested per socket. Only set if the --ntasks-per-socket option is specified.
SLURM_SUBMIT_DIR	PBS_O_WORKDIR	The directory from which sbatch was invoked
SLURM_SUBMIT_HOST	PBS_O_HOST	The hostname of the computer from which sbatch was invoked
SLURM_TASK_PID		The process ID of the task being started
SLURMD_NODENAME		Name of the node running the job script
SLURM_JOB_GPUS		GPU IDs allocated to the job (if any).

Does not work on AWS!

<https://wiki.hpcc.msu.edu/display/ITH/SLURM+Environment+Variables>

# Optimizing Parallel Processing

## Parallel Processing (5 jobs)

Increasing memory per job

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	ODELIST (REASON)	
2	queue1	rplp1	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	← 64 GB RAM
3	queue1	rs1	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
4	queue1	rs2	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
5	queue1	rs3	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	
6	queue1	rs4	ec2-user	R	0:02	1	queue1-dy-r52xlarge-1	

↑  
40 GB RAM

```
#SBATCH --time=24:00:00 # walltime
#Sbatch --ntasks=1 # number of processor cores (i.e. tasks)
#SBATCH --nodes=1 # number of nodes
#SBATCH --cpus-per-task=5 # number of processors per task
#SBATCH -J "rs1" # job name
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	ODELIST (REASON)
7	queue1	rplp1	ec2-user	R	4:31	1	queue1-dy-r52xlarge-1
8	queue1	rs1	ec2-user	R	4:31	1	queue1-dy-r52xlarge-1
9	queue1	rs2	ec2-user	R	1:25	1	queue1-dy-r52xlarge-2
10	queue1	rs3	ec2-user	R	1:25	1	queue1-dy-r52xlarge-3
11	queue1	rs4	ec2-user	R	1:25	1	queue1-dy-r52xlarge-4

# Parallel Data Processing

## Parallel Processing

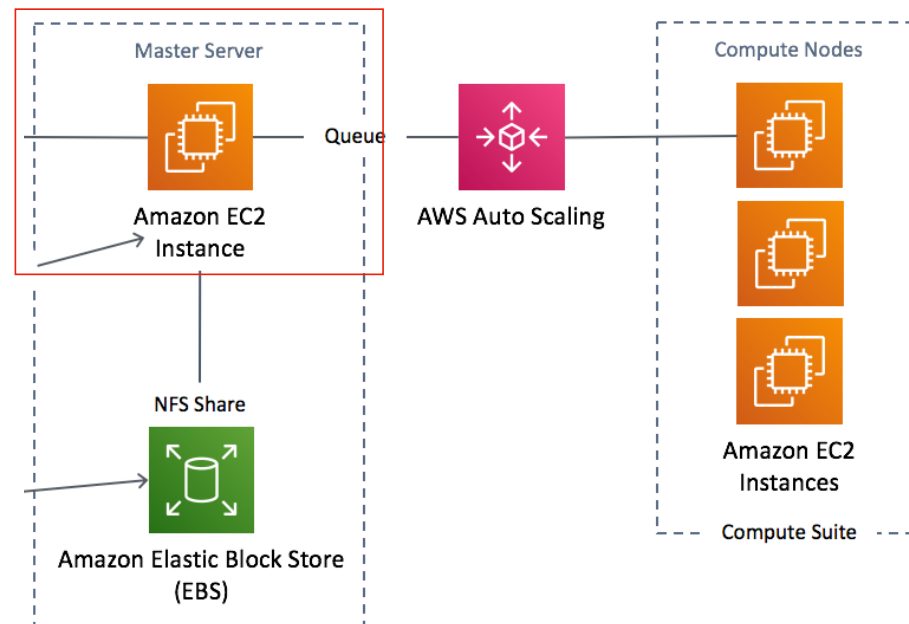
- Lab 4 techniques
- Parallelizing jobs
- Increasing memory per job
- Terminating cluster head nodes automatically
- Slurm job dependencies



# AWS ParallelCluster

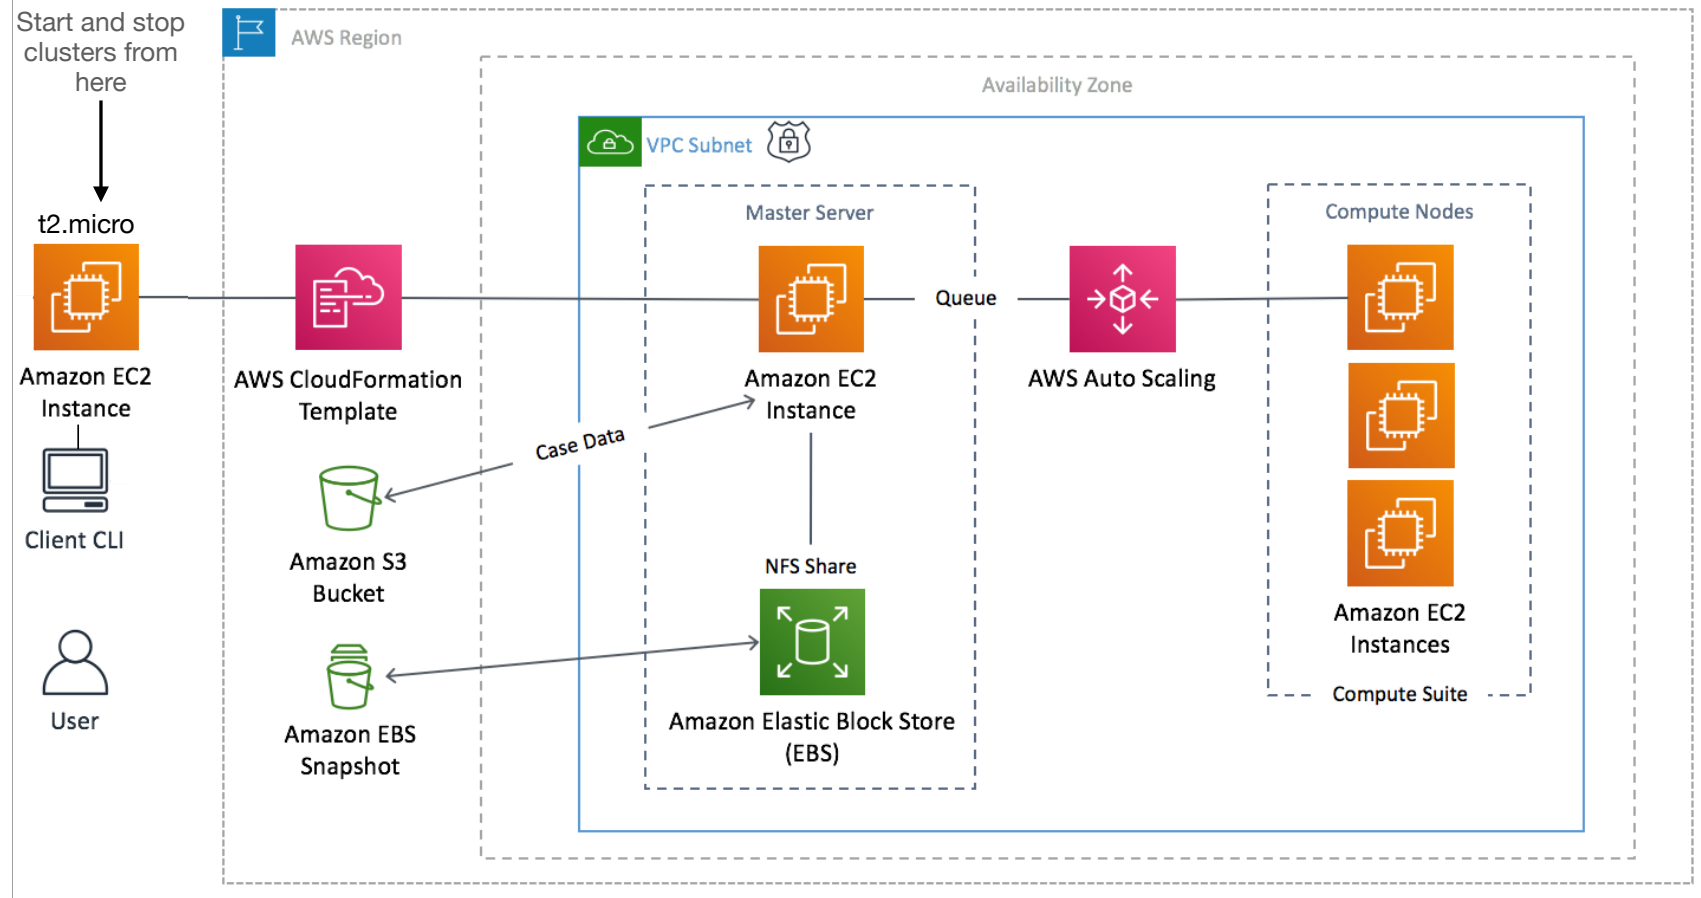
## Master Server

- Master node
- Main login interface to the cluster
- Jobs are submitted and managed here
- Always running once the cluster is created
- Configured using the file `~/cluster-config.yaml`



# Data Processing on AWS

## Cluster Deletion



# Data Processing on AWS

## Cluster Deletion

Top Free Tier Services by Usage			<a href="#">View all</a>
Service	Free Tier usage limit	Month-to-date usage	
Amazon Elastic Compute Cloud	1 GB of Amazon Elastic Block Storage snapshot storage	<b>100.00%</b> (1.00/1 GB-mo)	
Amazon Elastic Compute Cloud	30 GB of Amazon Elastic Block Storage in any combination of General Purpose (SSD) or Magnetic	<b>100.00%</b> (30.00/30 GB-Mo)	
Amazon Elastic Compute Cloud	750 hours of Amazon EC2 Linux t2.micro instance usage	27.33% (205.00/750 Hrs)	750/24 = 31.25 days
Amazon Simple Storage Service	2,000 Put, Copy, Post or List Requests of Amazon S3	3.30% (66.00/2,000 Requests)	
AmazonCloudWatch	5 GB of Log Data Ingestion for Amazon Cloudwatch	0.97% (0.05/5 GB)	

# Data Processing on AWS

## Cluster Deletion

- Normal procedure:
  - receive notification that jobs have been completed
  - `update_snapshot.sh data 2`
  - receive snapshot completion notification from CloudWatch
  - `pcluster delete-cluster -n MyCluster01`

# Data Processing on AWS

## Cluster Deletion

- Could have EC2 instance constantly check on queue in Head Node to see if it is empty
  - `pcluster ssh -i ~/MyKeyPair.pem -n MyCluster01 "queue"`
  - and then create snapshot and delete cluster
  - Very inefficient

# Data Processing on AWS

## Cluster Deletion

- Can use slurm job dependencies to ssh to EC2 instance to issue `pcluster delete-cluster` command

```
PyHipp — ec2-user@ip-10-0-7-189:/data/src/PyHipp — -zsh — 93x24
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch sleepslurm.sh
Submitted batch job 38
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch --dependency=afterok:38 /data/src/PyHipp/consol
l_jobs.sh
Submitted batch job 39
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ squeue
      JOBID PARTITION    NAME   USER  ST       TIME  NODES NODELIST(REASON)
       39   queue1  consol_j  ec2-user  PD        0:00      1 (Dependency)
       38   queue1  example-  ec2-user   R        0:12      1 queue1-dy-r5a2xlarge-1
```

- `ssh -i /data/MyKeyPair.pem ec2-user@xx.xx.xx.xx “~/update_snapshot.sh data 2 MyCluster01; pcluster delete-cluster -n MyCluster01”`
- but we want to delete the cluster only after the snapshot is complete

# Data Processing on AWS

## Cluster Deletion

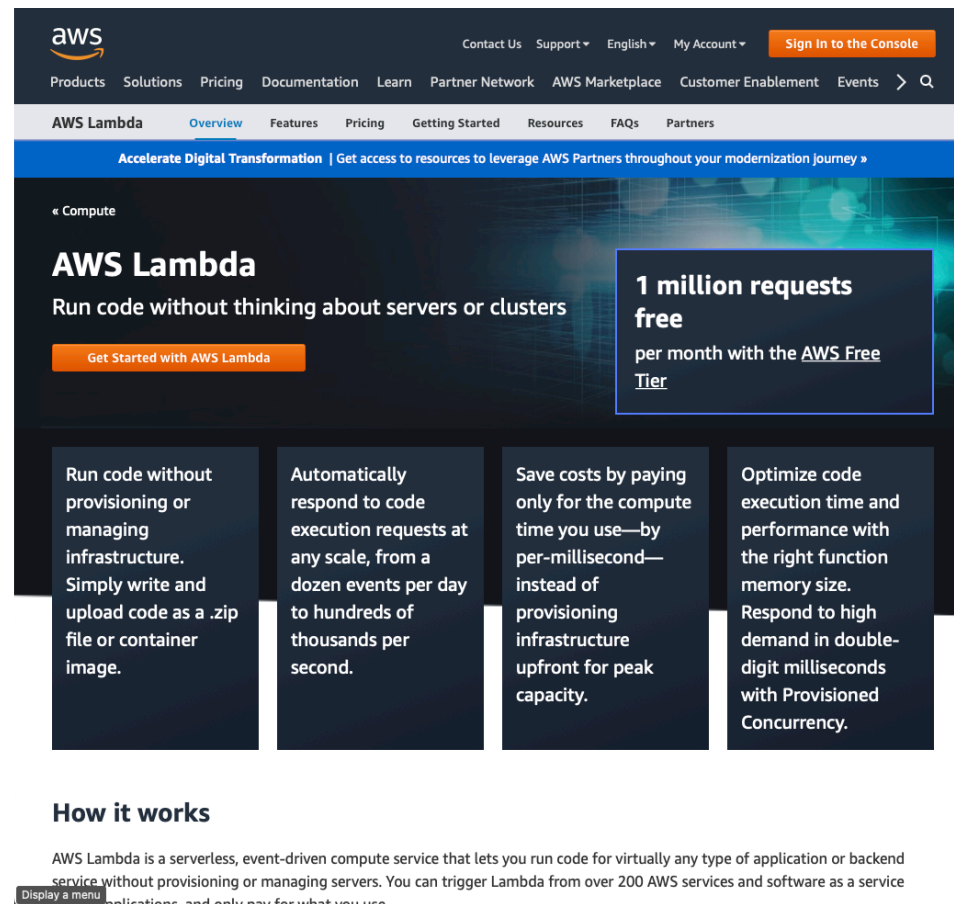
Can we use CloudWatch/EventBridge instead?

The screenshot shows the AWS Management Console interface. The left sidebar contains the navigation menu with 'CloudWatch' selected. The main content area features a notification banner stating 'CloudWatch Events is now Amazon EventBridge'. Below this, the 'Rules' section is visible, including a 'Create rule' button and a table of existing rules.

Status	Name	Description
<input checked="" type="radio"/>	EC2Running	
<input checked="" type="radio"/>	SnapshotComplete	

# Data Processing on AWS

## Cluster Deletion



The screenshot shows the AWS Lambda landing page. At the top is the AWS logo and navigation links: Contact Us, Support, English, My Account, and a Sign In to the Console button. Below this is a secondary navigation bar with links for Products, Solutions, Pricing, Documentation, Learn, Partner Network, AWS Marketplace, Customer Enablement, and Events. The main header for AWS Lambda includes links for Overview, Features, Pricing, Getting Started, Resources, FAQs, and Partners. A blue banner below the header reads "Accelerate Digital Transformation | Get access to resources to leverage AWS Partners throughout your modernization journey". The main content area features the heading "AWS Lambda" and the subtext "Run code without thinking about servers or clusters". A prominent orange button says "Get Started with AWS Lambda". To the right, a dark box highlights "1 million requests free per month with the AWS Free Tier". Below this are four columns of text describing the service's benefits: running code without provisioning infrastructure, automatic response to code execution requests at any scale, saving costs by paying only for compute time, and optimizing code execution time and performance with the right function memory size. At the bottom, a section titled "How it works" begins with the text: "AWS Lambda is a serverless, event-driven compute service that lets you run code for virtually any type of application or backend service without provisioning or managing servers. You can trigger Lambda from over 200 AWS services and software as a service applications, and only pay for what you use."

« Compute

## AWS Lambda

Run code without thinking about servers or clusters

[Get Started with AWS Lambda](#)

**1 million requests free**  
per month with the [AWS Free Tier](#)

- Run code without provisioning or managing infrastructure. Simply write and upload code as a .zip file or container image.
- Automatically respond to code execution requests at any scale, from a dozen events per day to hundreds of thousands per second.
- Save costs by paying only for the compute time you use—by per-millisecond—instead of provisioning infrastructure upfront for peak capacity.
- Optimize code execution time and performance with the right function memory size. Respond to high demand in double-digit milliseconds with Provisioned Concurrency.

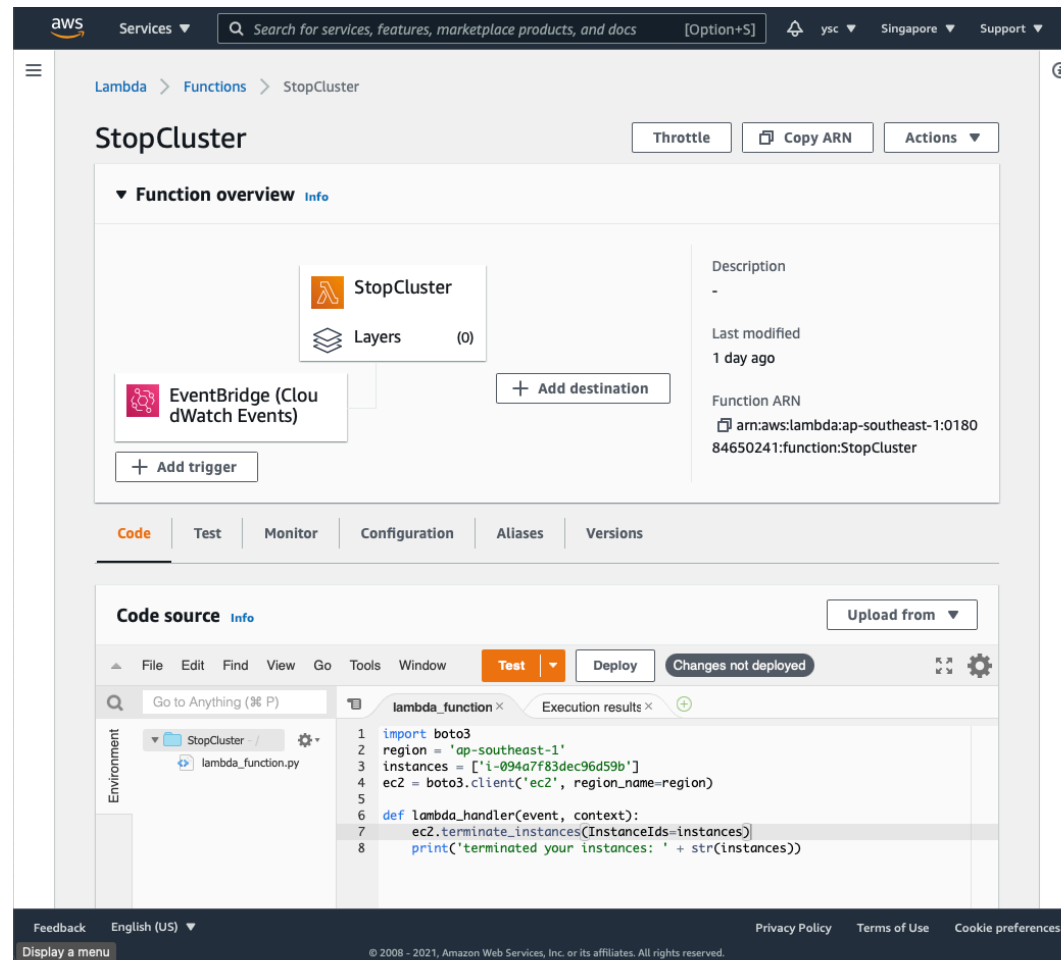
### How it works

AWS Lambda is a serverless, event-driven compute service that lets you run code for virtually any type of application or backend service without provisioning or managing servers. You can trigger Lambda from over 200 AWS services and software as a service applications, and only pay for what you use.



# Data Processing on AWS

## Cluster Deletion



# Data Processing on AWS

## Cluster Deletion

- Use slurm job dependencies to initiate snapshot:

```
PyHipp — ec2-user@ip-10-0-7-189:/data/src/PyHipp — zsh — 93x24
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch sleepslurm.sh
Submitted batch job 38
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch --dependency=afterok:38 /data/src/PyHipp/consol_jobs.sh
Submitted batch job 39
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ squeue
      JOBID PARTITION   NAME   USER  ST       TIME  NODES NODELIST(REASON)
       39   queue1  consol_j ec2-user PD        0:00      1 (Dependency)
       38   queue1  example- ec2-user  R        0:12      1 queue1-dy-r5a2xlarge-1
```

- `ssh -i /data/MyKeyPair.pem ec2-user@xx.xx.xx.xx "~/update_snapshot.sh data 2 MyCluster01"`
- CloudWatch watch for snapshot completion
- Lambda function used to delete head node

# Data Processing on AWS

## Cluster Deletion

- Use slurm job dependencies to initiate snapshot:

```
PyHipp — ec2-user@ip-10-0-7-189:/data/src/PyHipp — zsh — 93x24
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch sleepslurm.sh
Submitted batch job 38
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ sbatch --dependency=afterok:38 /data/src/PyHipp/consol_jobs.sh
Submitted batch job 39
(env1) [ec2-user@ip-10-0-8-173 PyHipp]$ squeue
      JOBID PARTITION    NAME   USER  ST       TIME  NODES NODELIST(REASON)
       39   queue1  consol_j  ec2-user  PD      0:00      1 (Dependency)
       38   queue1  example-  ec2-user  R      0:12      1 queue1-dy-r5a2xlarge-1
```

- `ssh -o StrictHostKeyChecking=no -i /data/MyKeyPair.pem ec2-user@xx.xx.xx.xx "source ~/.bash_profile; pcluster update-compute-fleet --status STOP_REQUESTED -n MyCluster01; ~/update_snapshot.sh data 2 MyCluster01"`
- CloudWatch watch for snapshot completion
- Lambda function used to delete head node
- Still have to call `pcluster delete-cluster` to remove from list

# Parallel Data Processing

## Parallel Processing

- Lab 4 techniques
- Parallelizing jobs
- Increasing memory per job
- Terminating cluster head nodes automatically
- Slurm job dependencies

# Data Processing on AWS

## Cluster Deletion

### consol\_jobs.sh

```
#!/bin/sh
templ=($(squeue))
cmd1="sbatch --dependency=afterok:"
counter1=0
for i in "${templ[@]"; do
    if [[ "$i" == "queue1" ]]; then
        id1=${templ[$counter1-1]}
        cmd1="${cmd1}${id1}:"
    fi
    counter1=$((counter1+1))
done
cmd1=${cmd1::-1}
cmd1="${cmd1} /data/src/PyHipp/ec2snapshot.sh"

echo $cmd1
eval $cmd1
```

```
(env1) [ec2-user@ip-10-0-3-231 20181105]$ squeue
  JOBID PARTITION   NAME   USER ST   TIME  NODES NODELIST(REASON)
   233   queue1    rplpl ec2-user CF    0:02    1 compute-dy-r5large-1
   234   queue1      rse ec2-user CF    0:02    1 compute-dy-r5large-1
(env1) [ec2-user@ip-10-0-3-231 20181105]$
```

# Data Processing on AWS

## Cluster Deletion

consol\_jobs.sh

```
#!/bin/sh
templ=$(squeue)
cmd1="sbatch --dependency=afterok:"
counter1=0
for i in "${templ[@]}"; do
    if [[ "$i" == "queue1" ]]; then
        id1=${templ[$counter1-1]}
        cmd1="${cmd1}${id1}:"
    fi
    counter1=$((counter1+1))
done
cmd1=${cmd1::-1}
cmd1="${cmd1} /data/src/PyHipp/ec2snapshot.sh"
echo ${cmd1}
eval $cmd1
```

```
(env1) [ec2-user@ip-10-0-9-92 data]$ for i in "${templ[@]}"; do echo $i; done
JOBID
PARTITION
NAME
USER
ST
TIME
NODES
NODELIST(REASON)
233
queue1
rplpl
ec2-user
CF
0:02
1
compute-dy-r5large-1
234
queue1
rse
ec2-user
CF
0:02
1
compute-dy-r5large-1
(env1) [ec2-user@ip-10-0-9-92 data]$
```

```
cmd1="sbatch --dependency=afterok:
cmd1="sbatch --dependency=afterok:233:
cmd1="sbatch --dependency=afterok:233:234:
cmd1="sbatch --dependency=afterok:233:234
cmd1="sbatch --dependency=afterok:233:234 /data/src/PyHipp/ec2snapshot.sh
```

# Optimizing Parallel Processing

## Lab 6 Techniques

- Bash
  - GitHub command line (git pull, git push, git fetch upstream, etc.)
  - Create shell scripts for frequently performed series of functions
  - Run commands remotely via ssh, e.g.

```
pcluster ssh -i ~/MyKeyPair.pem -n MyCluster01 squeue
```
- Slurm
  - Use file dependencies to parallelize into 2 jobs (RPLParallel, RPLSplit)
  - Use directory hierarchy to parallelize into 5 jobs (e.g. session01/array01)
  - Use slurm parameter --cpus-per-task=5 to increase memory available to individual jobs
  - Use slurm job dependencies to run jobs in sequence
- AWS
  - Use AWS Lambda to run scripts without servers
  - Use AWS Lambda, EventBridge, and small EC2 instance to terminate head node automatically

# Optimizing Parallel Processing

## Lab Instructions

- Lab 6 Instructions:
  - <https://ee3801.github.io/Lab6/instruction.html>
- Submit to Canvas (Lab 6->Lab 6A & Lab 6->Lab 6B)
- Submit in PDF format
- Name the files Lab6A\_YourName.pdf and Lab6B\_YourName.pdf
- Part A due on Monday (Oct 23) 2 pm
- Part B due on Wednesday (Oct 25) 9 pm



# Questions?